

2. Dataset Selection and Data Extraction

```
In [11]: import pandas as pd

data = pd.read_csv('HRDataSets.csv')

print("Dataset Info:")
data.info()

print("\n Dataset shape:", data.shape)

print("\n Value counts of Attrition:")
print(data['Attrition'].value_counts())

print("\n First 5 rows of the dataset:")
data.head()
```

Dataset Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 11906 entries, 0 to 11905

Data columns (total 23 columns):

#	Column	Non-Null Count	Dtype
0	Employee ID	11906 non-null	int64
1	Age	11906 non-null	int64
2	Gender	11906 non-null	object
3	Years at Company	11906 non-null	int64
4	Job Role	11906 non-null	object
5	Monthly Income	11906 non-null	int64
6	Work-Life Balance	11906 non-null	object
7	Job Satisfaction	11906 non-null	object
8	Performance Rating	11906 non-null	object
9	Number of Promotions	11906 non-null	int64
10	Overtime	11906 non-null	object
11	Distance from Home	11906 non-null	int64
12	Education Level	11906 non-null	object
13	Marital Status	11906 non-null	object
14	Company Tenure	11906 non-null	int64
15	Number of Dependents	11906 non-null	int64
16	Job Level	11906 non-null	object
17	Remote Work	11906 non-null	object
18	Leadership Opportunities	11906 non-null	object
19	Innovation Opportunities	11906 non-null	object
20	Company Reputation	11906 non-null	object
21	Employee Recognition	11906 non-null	object
22	Attrition	11906 non-null	object

dtypes: int64(8), object(15)

memory usage: 2.1+ MB

Dataset shape: (11906, 23)

Value counts of Attrition:

Attrition

Stayed 6278

Left 5628

Name: count, dtype: int64

First 5 rows of the dataset:

Out[11]:

	Employee ID	Age	Gender	Years at Company	Job Role	Monthly Income	Work-Life Balance	Job Satisfaction	Performance Rating
0	30585	35	Male	7	Education	4563	Good	High	Average
1	54656	50	Male	7	Education	5583	Fair	High	Average
2	33442	58	Male	44	Media	5525	Fair	Very High	
3	46775	22	Female	5	Healthcare	8700	Good	High	Average
4	65181	55	Female	16	Media	5939	Poor	High	Average

5 rows × 23 columns



3. Data Cleaning and Preprocessing

```
In [13]: print("\nMissing values per column:")
print(data.isnull().sum())
```

```
Missing values per column:
Employee ID          0
Age                  0
Gender               0
Years at Company     0
Job Role             0
Monthly Income       0
Work-Life Balance    0
Job Satisfaction     0
Performance Rating   0
Number of Promotions 0
Overtime             0
Distance from Home   0
Education Level      0
Marital Status       0
Number of Dependents 0
Job Level            0
Remote Work          0
Leadership Opportunities 0
Innovation Opportunities 0
Company Reputation   0
Employee Recognition 0
Attrition            0
dtype: int64
```

```
In [15]: data.describe().round(0).astype(int)
```

Out[15]:

	Employee ID	Age	Years at Company	Monthly Income	Number of Promotions	Distance from Home	Number of Dependents
count	11906	11906	11906	11906	11906	11906	11906
mean	37151	38	16	7298	1	50	2
std	21479	12	11	2154	1	29	2
min	5	18	1	1226	0	1	0
25%	18580	28	7	5645	0	25	0
50%	37005	38	13	7346	1	50	1
75%	55732	49	22	8862	2	75	3
max	74465	59	51	15063	4	99	6

In [17]: `print("\nData types of each column:")`
`data.dtypes`

Data types of each column:

Out[17]: Employee ID int64
 Age int64
 Gender object
 Years at Company int64
 Job Role object
 Monthly Income int64
 Work-Life Balance object
 Job Satisfaction object
 Performance Rating object
 Number of Promotions int64
 Overtime object
 Distance from Home int64
 Education Level object
 Marital Status object
 Number of Dependents int64
 Job Level object
 Remote Work object
 Leadership Opportunities object
 Innovation Opportunities object
 Company Reputation object
 Employee Recognition object
 Attrition object
 dtype: object

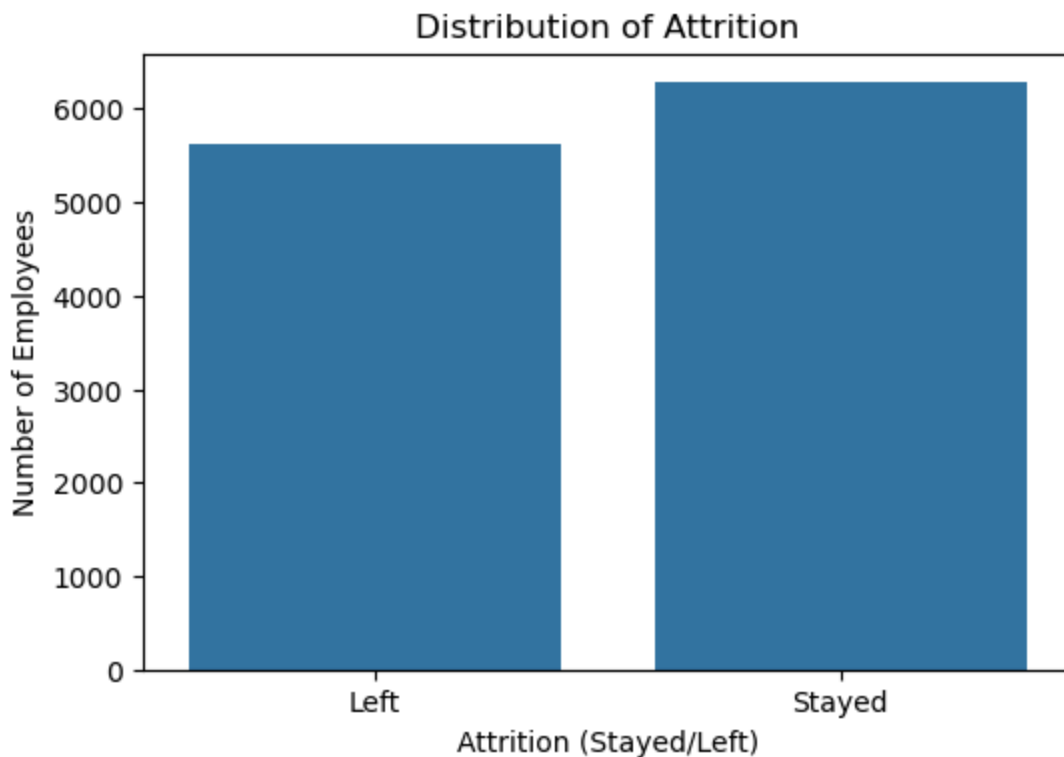
In [12]: `if 'Company Tenure' in data.columns:`
`data = data.drop('Company Tenure', axis=1)`
`print("\n'Company Tenure' column removed.")`
`else:`
`print("\n'Company Tenure' column not found.")`

'Company Tenure' column removed.

4. Exploratory Data Analytics (EDA)

```
In [19]: import matplotlib.pyplot as mp
import seaborn as sb

mp.figure(figsize=(6, 4))
sb.countplot(data=data, x='Attrition')
mp.title('Distribution of Attrition')
mp.xlabel('Attrition (Stayed/Left)')
mp.ylabel('Number of Employees')
mp.show()
```





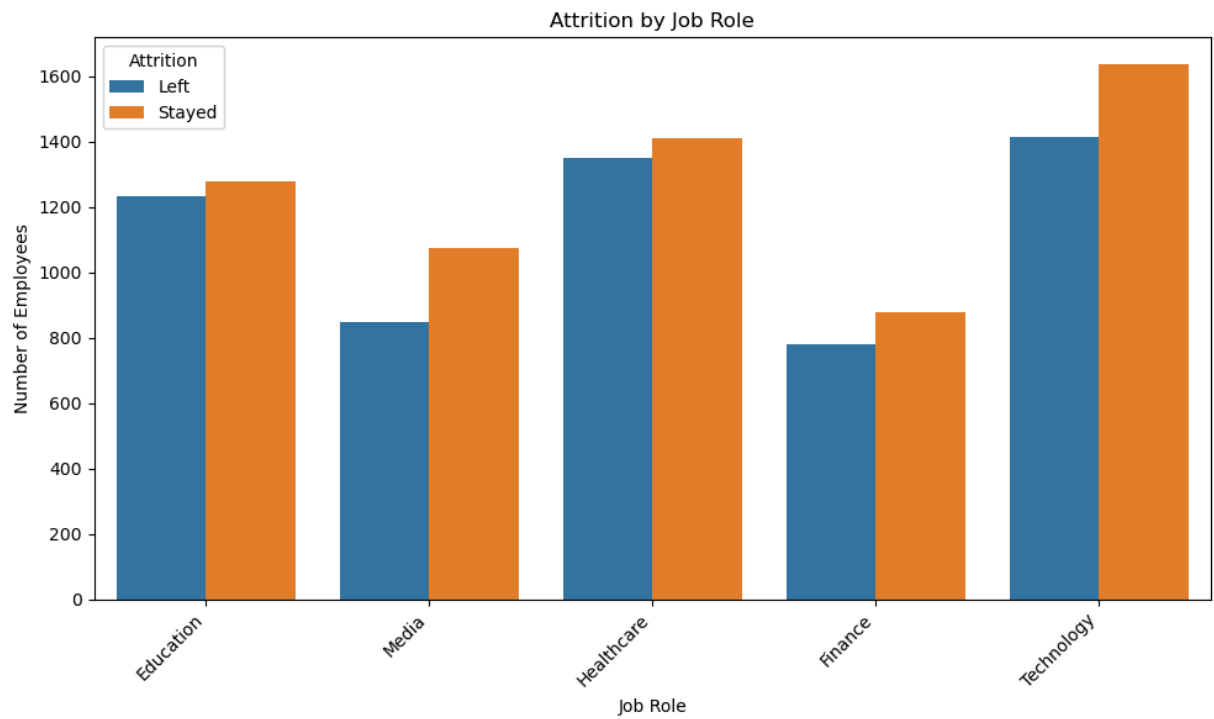
```
In [20]: # Attrition by Gender
mp.figure(figsize=(6, 4))
sb.countplot(data=data, x='Gender', hue='Attrition')
mp.title('Attrition by Gender')
mp.xlabel('Gender')
mp.ylabel('Number of Employees')
mp.show()


# Attrition by Job Role
mp.figure(figsize=(10, 6))
sb.countplot(data=data, x='Job Role', hue='Attrition')
mp.title('Attrition by Job Role')
mp.xlabel('Job Role')
mp.ylabel('Number of Employees')
mp.xticks(rotation=45, ha='right')
mp.tight_layout()
mp.show()
```

```
# Attrition by Work-Life Balance
mp.figure(figsize=(8, 5))
sb.countplot(data=data, x='Work-Life Balance', hue='Attrition', order=['Poor', 'Fair', 'Good'])
mp.title('Attrition by Work-Life Balance')
mp.xlabel('Work-Life Balance')
mp.ylabel('Number of Employees')
mp.show()


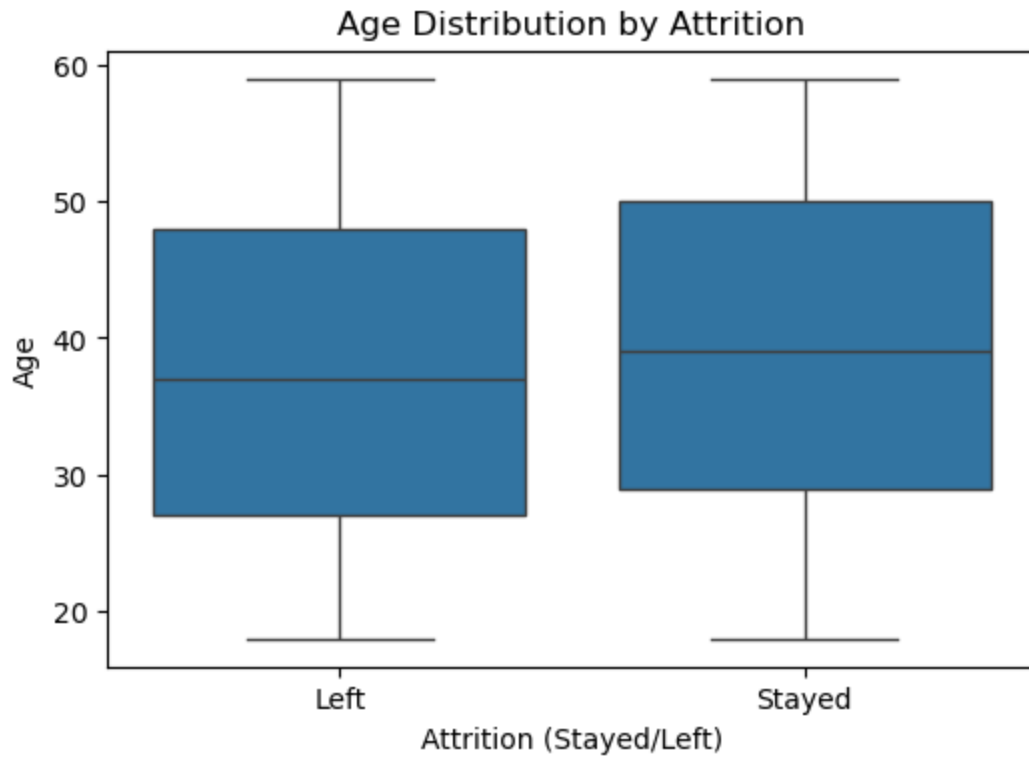
# Attrition by Job Satisfaction
mp.figure(figsize=(8, 5))
sb.countplot(data=data, x='Job Satisfaction', hue='Attrition', order=['Low', 'Medium', 'High'])
mp.title('Attrition by Job Satisfaction')
mp.xlabel('Job Satisfaction')
mp.ylabel('Number of Employees')
mp.show()
```

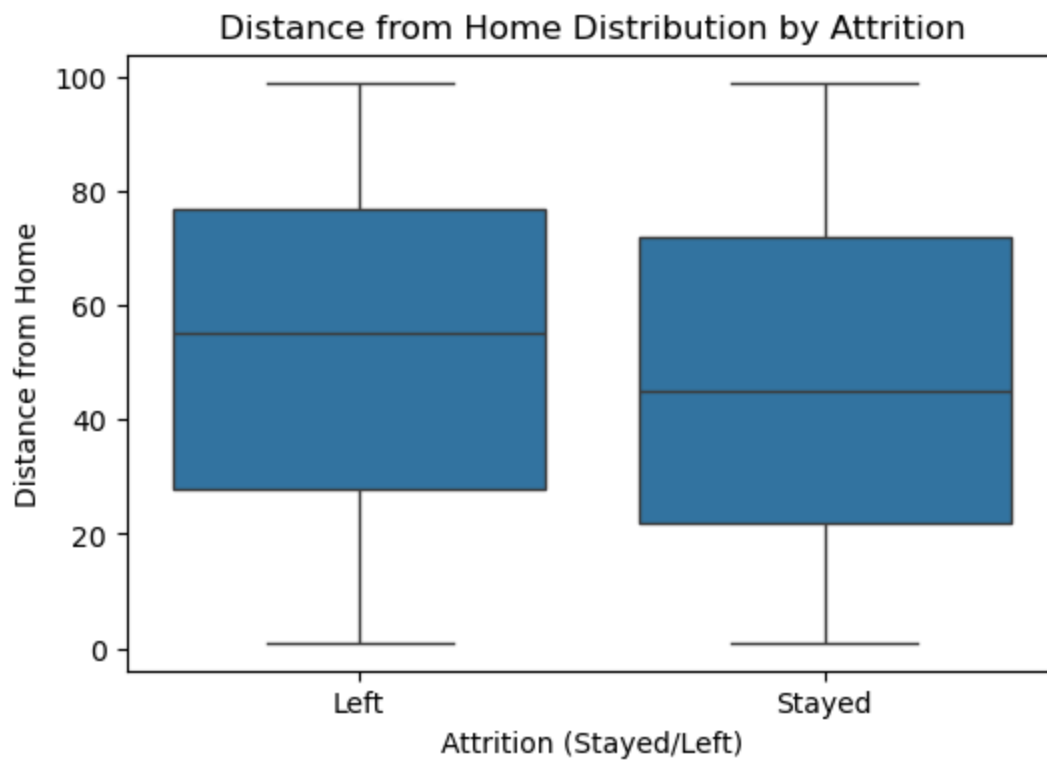




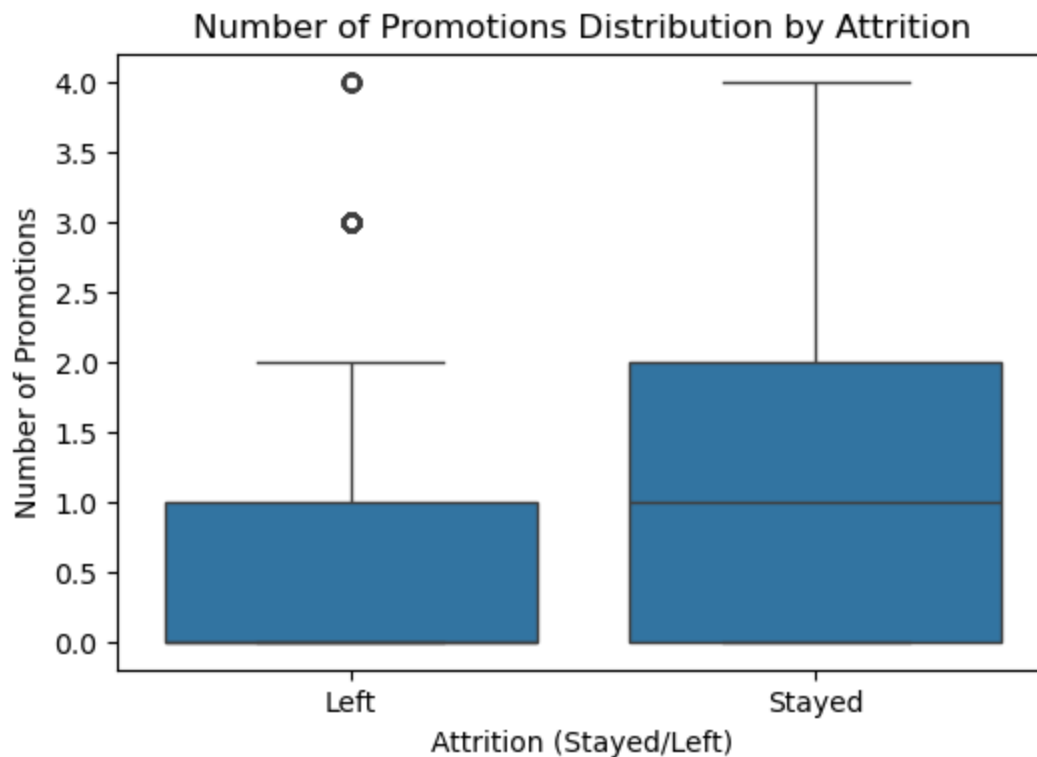




```
In [42]: numerical_columns = ['Age', 'Years at Company', 'Monthly Income', 'Distance from Ho
for col in numerical_columns:
    mp.figure(figsize=(4, ))
    sb.boxplot(data=data, x='Attrition', y=col)
    mp.title(f'{col} Distribution by Attrition')
    mp.xlabel('Attrition (Stayed/Left)')
    mp.ylabel(col)
    mp.show()
```






```
In [51]: numerical_data = data.select_dtypes(include=['number'])
cor_matrix = numerical_data.corr()
mp.figure(figsize=(10, 8))
sb.heatmap(cor_matrix, annot=True, cmap='RdBu', linewidths=0.5)
mp.title('Correlation Matrix of Numerical Variables')
mp.xticks(rotation = 45, ha='right')
mp.show()
```



```
In [52]: mp.figure(figsize=(8, 6))
sb.scatterplot(data=data, x='Age', y='Monthly Income', hue='Attrition')
mp.title('Scatter Plot of Age vs. Monthly Income by Attrition')
mp.xlabel('Age')
mp.ylabel('Monthly Income')
mp.show()
```



Project Implementation (Analytical Techniques and Methods - Implementation Focus)

Feature Encoding

```
In [54]: from sklearn.preprocessing import LabelEncoder

X = data.drop(['Employee ID', 'Attrition'], axis=1)
y = data['Attrition']

label_encoder = LabelEncoder()
y = label_encoder.fit_transform(y)

X = pd.get_dummies(X, drop_first=True)
print("\nFeatures (X) after one-hot encoding (first 5 rows):")
X.head()
```

Features (X) after one-hot encoding (first 5 rows):

Out[54]:

	Age	Years at Company	Monthly Income	Number of Promotions	Distance from Home	Number of Dependents	Gender_Male	Role_Financ
0	35	7	4563	1	55	4	True	Fals
1	50	7	5583	3	14	2	True	Fals
2	58	44	5525	0	43	4	True	Fals
3	22	5	8700	0	2	0	False	Fals
4	55	16	5939	0	31	1	False	Fals

5 rows × 38 columns



Data Splitting

```
In [55]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print("\nShape of X_train:", X_train.shape)
print("Shape of X_test:", X_test.shape)
print("Shape of y_train:", y_train.shape)
print("Shape of y_test:", y_test.shape)
```

Shape of X_train: (9524, 38)

Shape of X_test: (2382, 38)

Shape of y_train: (9524,)

Shape of y_test: (2382,)

Logistic Regression Implementation

```
In [68]: from sklearn.linear_model import LogisticRegression

lm = LogisticRegression(random_state=42, solver='liblinear')
lm.fit(X_train, y_train)
y_pred_logistic = lm.predict(X_test)
print("\nLogistic Regression Model Trained.")
```

Logistic Regression Model Trained.

```
In [71]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
import matplotlib.pyplot as plt

accuracy = accuracy_score(y_test, y_pred_logistic)
precision = precision_score(y_test, y_pred_logistic)
recall = recall_score(y_test, y_pred_logistic)
f1 = f1_score(y_test, y_pred_logistic)
roc_auc = roc_auc_score(y_test, lm.predict_proba(X_test)[:, 1])
```

```
print("Logistic Regression Performance:")
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision (Predicting 'Left'): {precision:.4f}")
print(f"Recall (Identifying 'Left'): {recall:.4f}")
print(f"F1-Score: {f1:.4f}")
print(f"AUC-ROC: {roc_auc:.4f}")

fpr, tpr, thresholds = roc_curve(y_test, lm.predict_proba(X_test)[: , 1])
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (AUC = {roc_auc:.4f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate (FPR)')
plt.ylabel('True Positive Rate (TPR)')
plt.title('Receiver Operating Characteristic (ROC) Curve - Logistic Regression')
plt.legend(loc='lower right')
plt.show()
```

Logistic Regression Performance:


Accuracy: 0.7443

Precision (Predicting 'Left'): 0.7483

Recall (Identifying 'Left'): 0.7763

F1-Score: 0.7620

AUC-ROC: 0.8291



Random Forest Implementation

In [74]: `from sklearn.ensemble import RandomForestClassifier`

```
rfm = RandomForestClassifier(random_state=42)
rfm.fit(X_train, y_train)
y_pred_rf = rfm.predict(X_test)
print("\nRandom Forest Model Trained.")
```

Random Forest Model Trained.

In [77]: `from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
import matplotlib.pyplot as plt
import numpy as np`

```
accuracy_rf = accuracy_score(y_test, y_pred_rf)
precision_rf = precision_score(y_test, y_pred_rf)
recall_rf = recall_score(y_test, y_pred_rf)
f1_rf = f1_score(y_test, y_pred_rf)
roc_auc_rf = roc_auc_score(y_test, rfm.predict_proba(X_test)[:, 1])
```

```
print("Random Forest Performance:")
print(f"Accuracy: {accuracy_rf:.4f}")
print(f"Precision (Predicting 'Left'): {precision_rf:.4f}")
print(f"Recall (Identifying 'Left'): {recall_rf:.4f}")
print(f"F1-Score: {f1_rf:.4f}")
print(f"AUC-ROC: {roc_auc_rf:.4f}")
```

```
fpr_rf, tpr_rf, thresholds_rf = roc_curve(y_test, rfm.predict_proba(X_test)[:, 1])
plt.figure(figsize=(8, 6))
plt.plot(fpr_rf, tpr_rf, color='darkgreen', lw=2, label=f'ROC curve (AUC = {roc_auc_rf:.4f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate (FPR)')
plt.ylabel('True Positive Rate (TPR)')
plt.title('Receiver Operating Characteristic (ROC) Curve - Random Forest')
plt.legend(loc='lower right')
plt.show()
```

```
importances = rfm.feature_importances_
feature_names = X_train.columns
sorted_indices = np.argsort(importances)[::-1]
```

```
plt.figure(figsize=(12, 8))
plt.title("Feature Importances - Random Forest")
plt.bar(range(X_train.shape[1]), importances[sorted_indices], align="center")
plt.xticks(range(X_train.shape[1]), feature_names[sorted_indices], rotation = 'vertical')
plt.tight_layout()
plt.show()
```


Random Forest Performance:

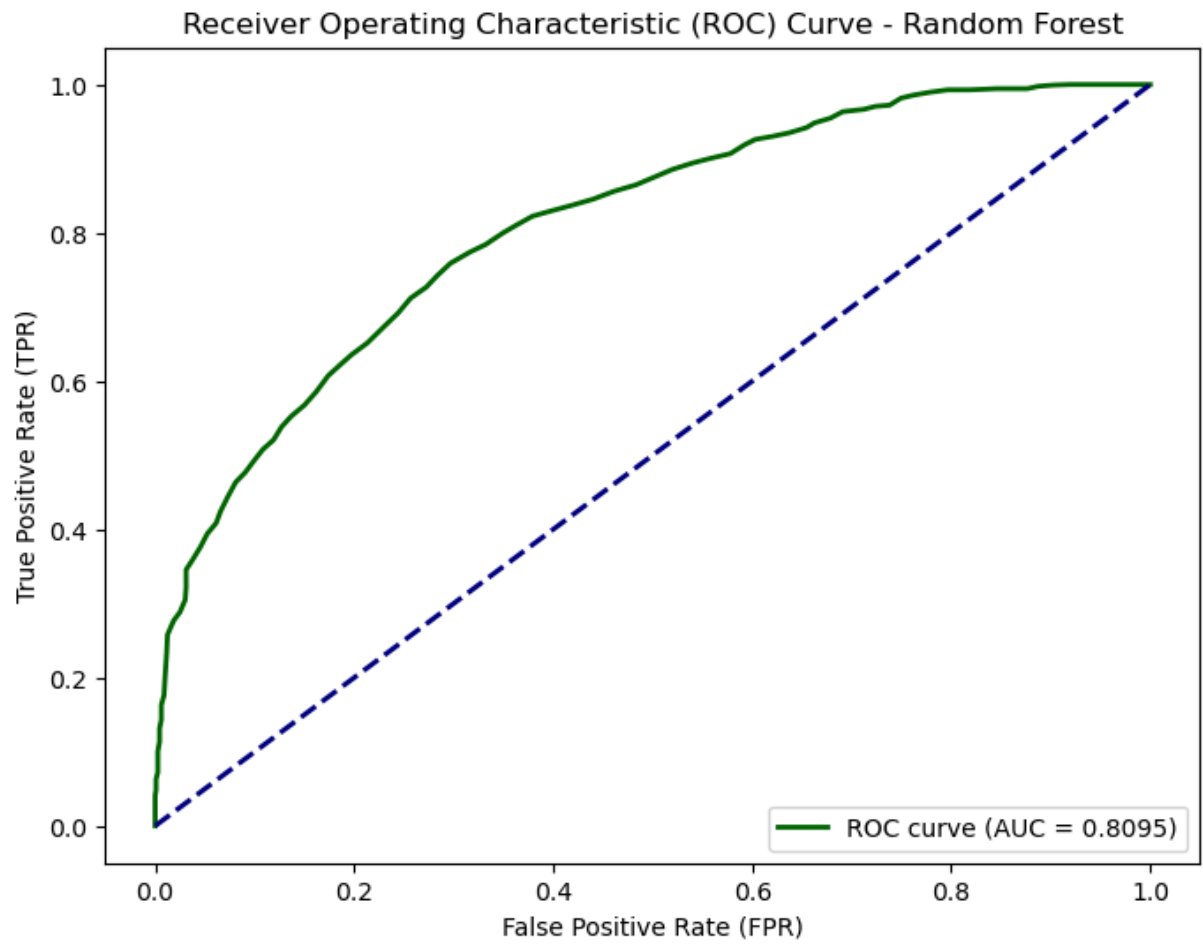
Accuracy: 0.7326


Precision (Predicting 'Left'): 0.7405

Recall (Identifying 'Left'): 0.7588

F1-Score: 0.7495

AUC-ROC: 0.8095





Data Analytics Artefact (Implemented Model)

```
In [62]: import joblib
joblib.dump(lm, 'attrition_model.pkl')
print("Logistic Regression model is saved to attrition_model.pkl")
```

Logistic Regression model is saved to attrition_model.pkl

```
In [64]: joblib.dump(rfm, 'attrition_model_rf.pkl')
print("Random Forest model is saved to attrition_model_rf.pkl")
```

Random Forest model is saved to attrition_model_rf.pkl

In []: